



# *Modeling VIII*

# Introduction



- Big Data
  - Large Sample Size
  - Large Number of Variables
  - Traditional Methods are Difficult to Implement
  - Depends on the Available Technology
- Goal: Explore Approaches for Quick Filtering of Predictors
- Tutorial 15
  - Download Rmd
  - Install Package `> library(glmnet)`
  - Knit the Document
  - Read the Introduction

## Introduction



My Data  
is Bigger than  
Your Data

## Linear Model



- Consider the Following:  
$$y_i = \beta_0 + X_{1i}\beta_1 + \dots + X_{pi}\beta_p + \epsilon_i$$
where  $i = 1, 2, 3, \dots, n$

- Matrix Representation

$$\mathbf{y} = \beta_0 + \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

$$\text{where } \mathbf{y} = [y_1, y_2, \dots, y_n]',$$

$$\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_p]',$$

$$\boldsymbol{\epsilon} = [\epsilon_1, \epsilon_2, \dots, \epsilon_n]',$$

and

$$\mathbf{X} = \begin{bmatrix} X_{11} & X_{21} & \dots & X_{p1} \\ X_{12} & X_{22} & \dots & X_{p2} \\ \vdots & \vdots & \ddots & \vdots \\ X_{1n} & X_{2n} & \dots & X_{pn} \end{bmatrix}$$

## Linear Model



- Information About Model Matrix

$$\mathbf{X} = \begin{bmatrix} X_{11} & X_{21} & \cdots & X_{p1} \\ X_{12} & X_{22} & \cdots & X_{p2} \\ \vdots & \vdots & \ddots & \vdots \\ X_{1n} & X_{2n} & \cdots & X_{pn} \end{bmatrix}$$

This Matrix Should Be Standardized

- Once Standardized, The Intercept  $\beta_0$  is Unnecessary in the Model
- For Interpretability, the Response Vector  $\mathbf{y}$  Can Also Be Standardized

## Part 1: Simulate and Meditate



- Run Chunk 1
  - Simulating Response From a Linear Model
  - All Predictor Variables in  $X$  are Standardized `> rnorm()`
  - What is  $n$ ?
  - What is  $p$ ?
  - What do We Know About the True Signal We Want to Detect?

# Sparse

## Part 1: Simulate and Meditate

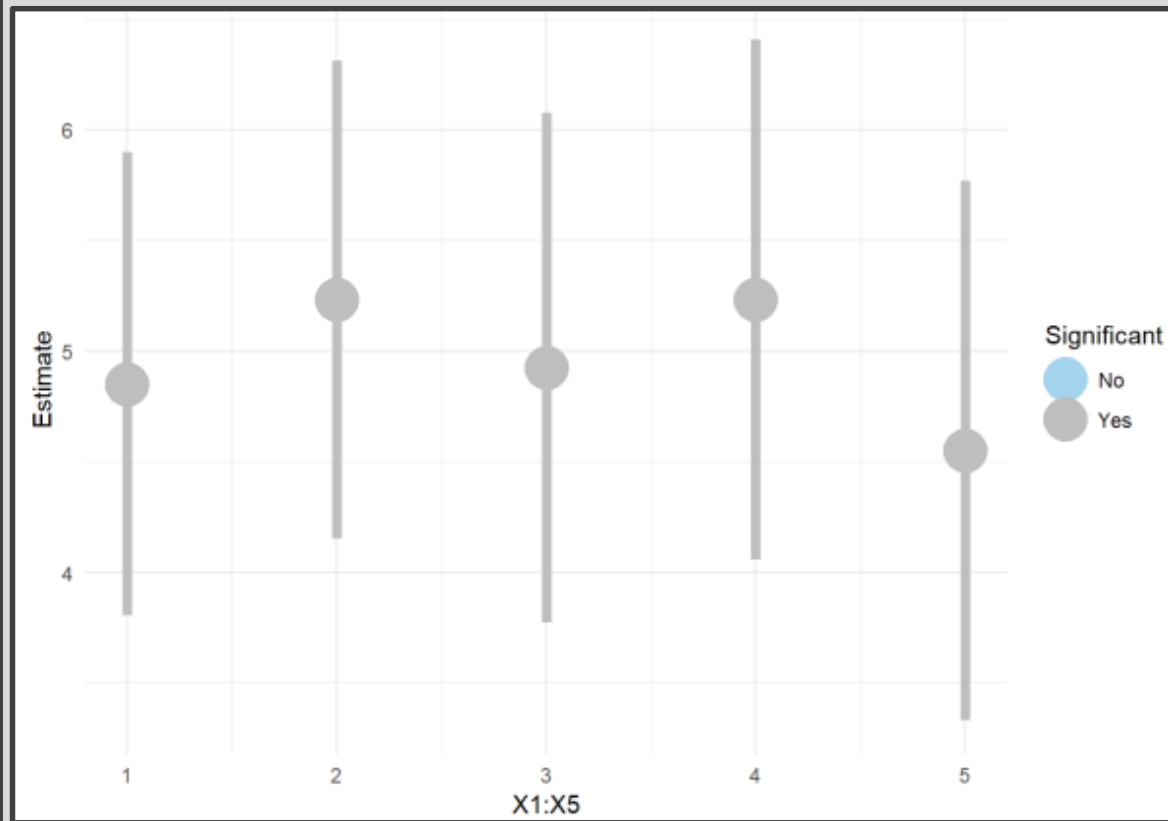


- Run Chunk 2
  - Fitting Naïve Linear Model
  - Obtaining Confidence Intervals for Parameters `> confint(lm.model)`
  - Figure Info
    - Show the Estimated Coefficients of Linear Model
    - Show Confidence Intervals for These Coefficients
    - What Does the Color Aesthetic Being Used For?

## Part 1: Simulate and Meditate



- Chunk 2 (Continued)
  - Knit the Document and Observe the 3 Graphics
  - Figure 1






## Part 1: Simulate and Meditate



- Chunk 2 (Continued)
  - Figure 2

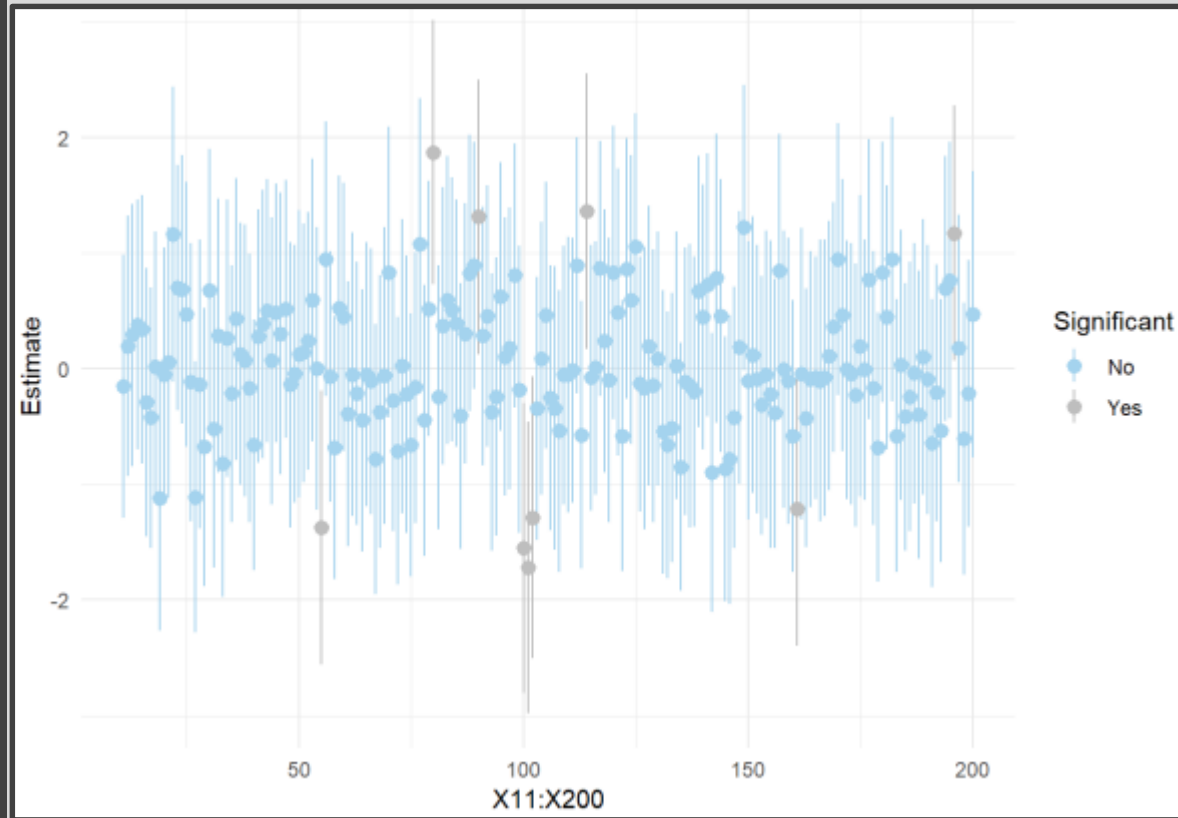


- What is the Problem?

## Part 1: Simulate and Meditate



- Chunk 2 (Continued)
  - Figure 3



- What is the Problem?

## Part 1: Simulate and Meditate



- Run Chunk 3
  - Regression for Each Predictor
  - Obtaining Coefficients

```
> coef(individual.mod)
(Intercept)      x.200
 0.1257668    -0.3200960
```

Save

- Obtaining P-Values

```
> summary(individual.mod)
```

Call:

```
lm(formula = y ~ ., data = SIM.DATA[, c(1, j + 1)])
```

Residuals:

Min	1Q	Median	3Q	Max
-47.252	-11.318	0.035	10.759	45.336

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.1258	0.7021	0.179	0.858
x.200	-0.3201	0.7230	-0.443	0.658

Save

Residual standard error: 15.66 on 498 degrees of freedom

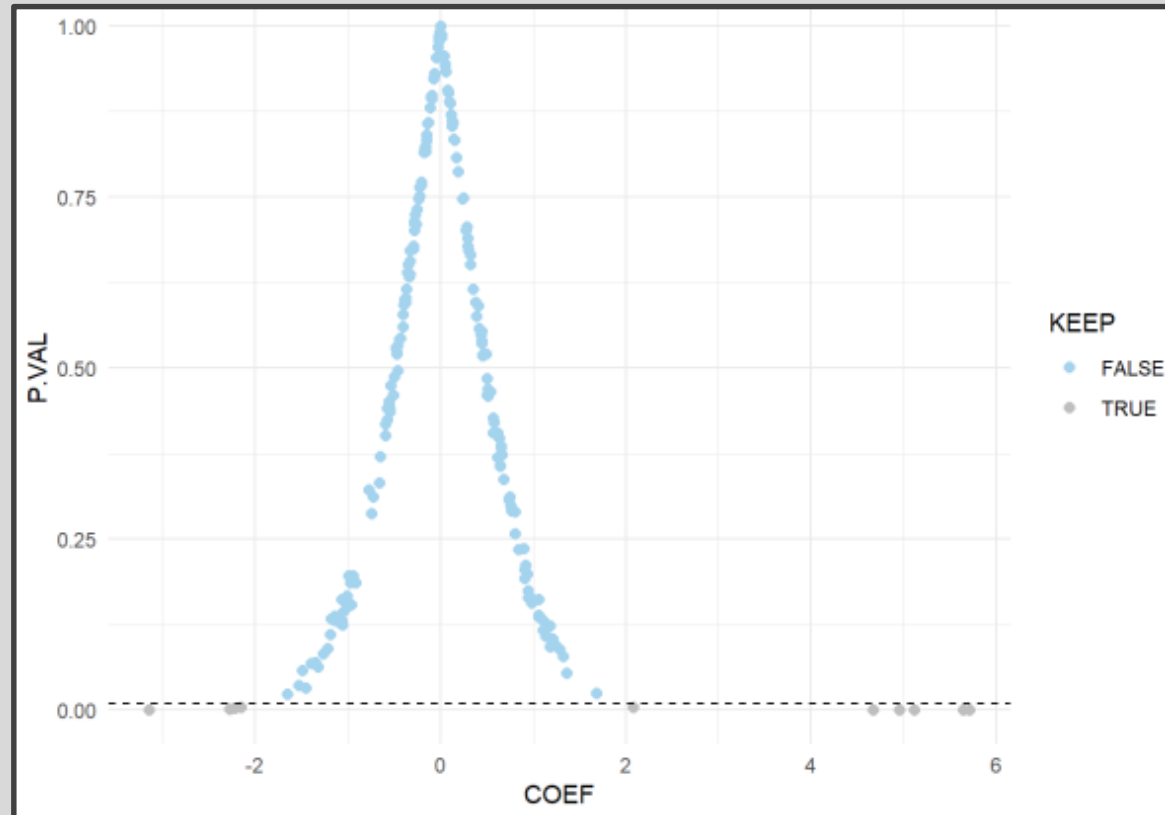
Multiple R-squared: 0.0003934, Adjusted R-squared: -0.001614

F-statistic: 0.196 on 1 and 498 DF, p-value: 0.6582

## Part 1: Simulate and Meditate



- Run Chunk 3
  - Figure Plots P-Values Against Coefficients



## Part 1: Simulate and Meditate



- Run Chunk 3
  - Suppose We Were to Keep Only the Predictor Variables that Had  $P\text{-Values} < 0.01$
  - Observe the Table

	P-Val > 0.01	P-Val < 0.01
Non-Zero	1%	4%
Zero	94%	1%

- 95% of Variables Ignored
- 5% of Variables Included
- Errors (What is Worse?)
  - We Will Ignore Variables that Are Important
  - We Will Include Variables that Are Irrelevant

## Part 1: Simulate and Meditate



- Chunk 4
  - Try to Find the Smallest Cutoff Value So That We are Not Missing Important Variables
  - To Ensure We are Not Missing Important Variables, Should we Increase or Decrease the Original Cutoff (0.01)
  - What Cutoff Works?
  - Try Multiple Cutoffs and Observe the Table
  - Run the Code Inside the Chunk Until All 10 Important Variables are Retained for the Future

## Part 1: Simulate and Meditate



- Chunk 4 (Continued)
  - Traditional Choice: 0.20
  - Output in Table

	P-Val > 0.01	P-Val < 0.01
Non-Zero	0%	5%
Zero	71%	24%

None of the Non-Zero Parameters Will Be Ignored

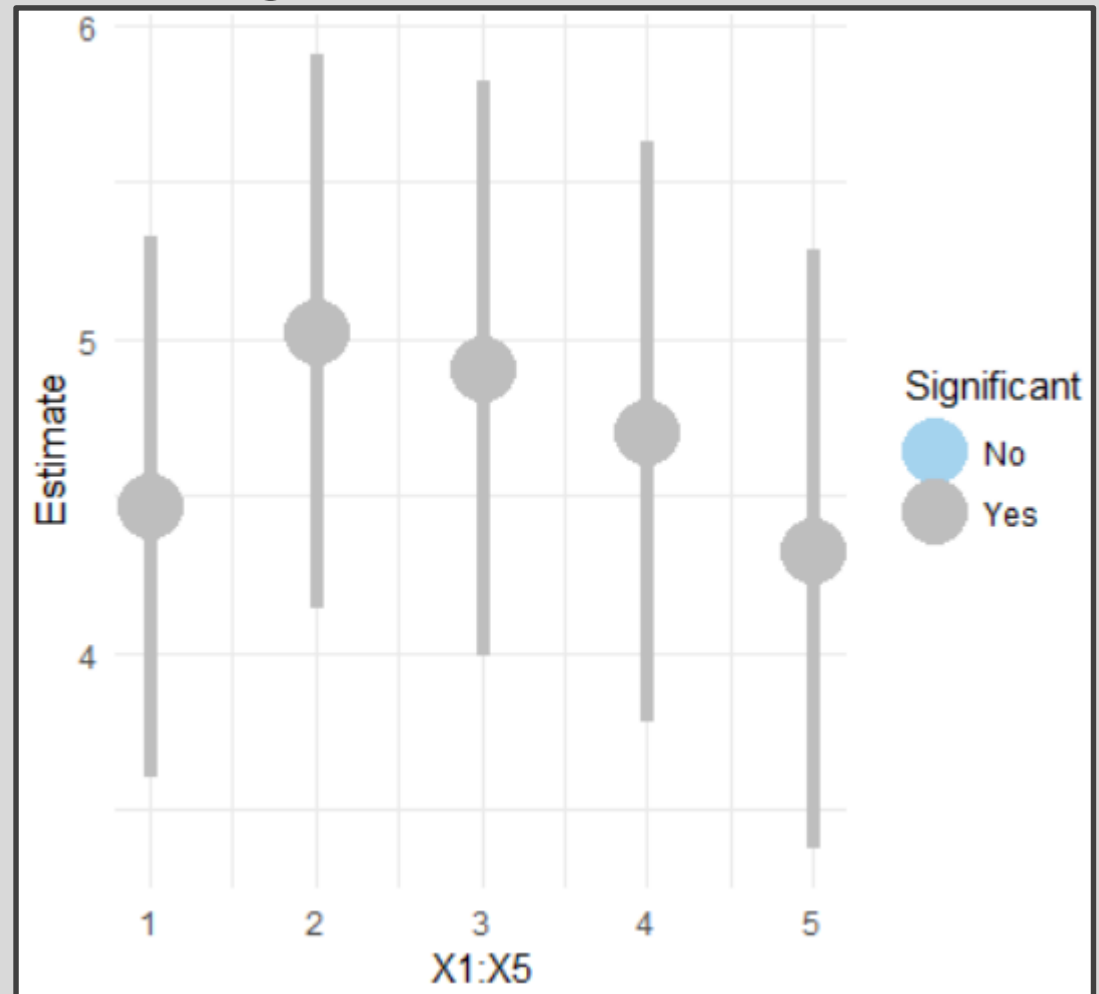
- Fit Linear Model for Variables Kept in Consideration

```
> lm(y~.,data=SIM.DATA[,c(1,which(KEEP)+1)])
```

## Part 2: Shrinkage Estimation



- Chunk 4 (Continued)
  - Suppose Cutoff is 0.2
  - Figure 1

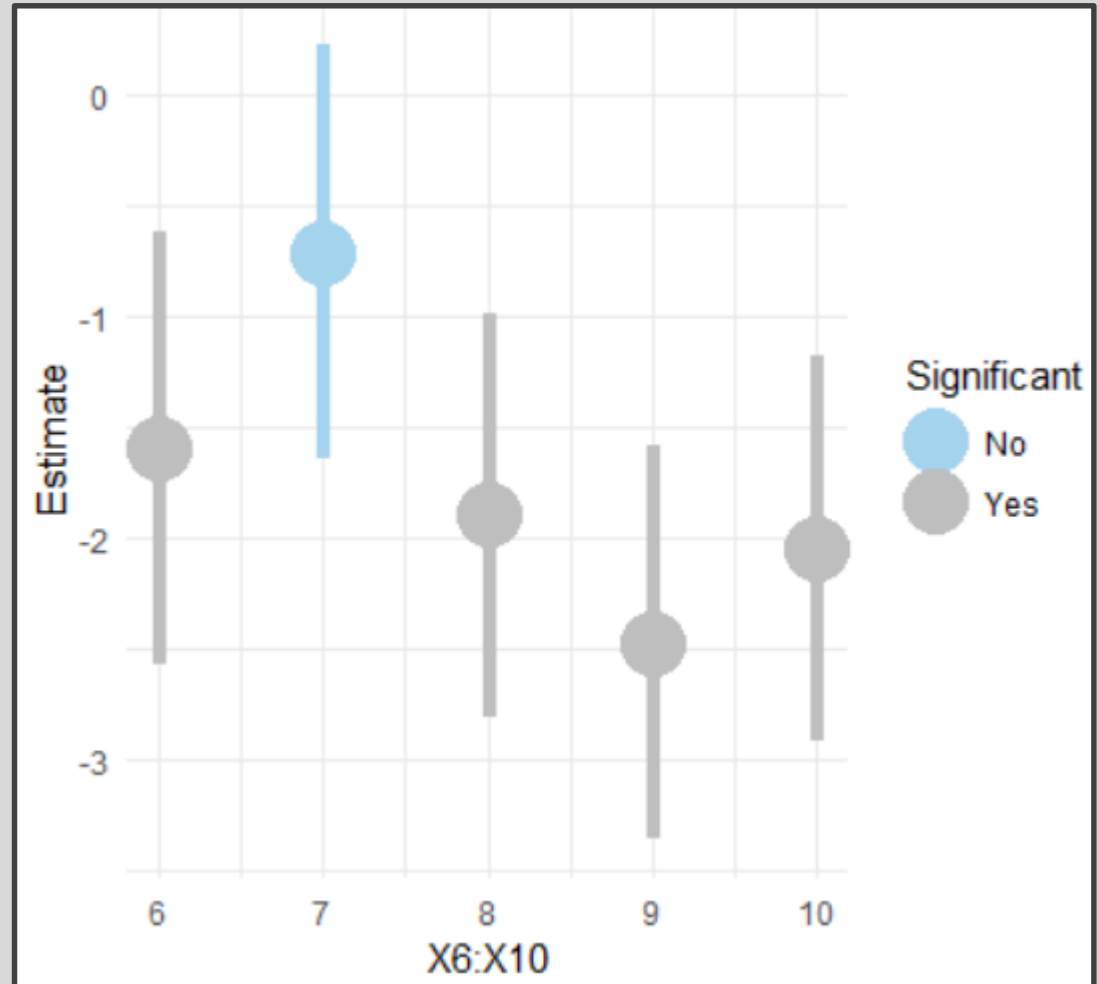




## Part 1: Simulate and Meditate




- Chunk 4 (Continued)
  - Figure 2



## Part 1: Simulate and Meditate



- Chunk 4 (Continued)
  - Figure 2



## Part 1: Simulate and Meditate



- Recap
  - Before Building Complex Models We are Performing a Simple Screening Procedure
  - Quick and Logical Approach
  - Problems
    - We May Lose Variables with Significant Interactions
    - We May Still Have Too Many
    - We May Retain Variables that are Highly Correlated
- Other Approach: Fit Full Model and Retain Variables with Sufficiently Small P-Values ( $<0.2$ )

## Part 2: Shrinkage Estimation and More Meditation



- Classic Linear Model Estimation
  - Minimize Sum of Squared Error
$$SSE = \sum [y_i - (\beta_0 + x_i' \boldsymbol{\beta})]^2$$
  - Optimization: Find  $\widehat{\beta}_0$  and  $\widehat{\boldsymbol{\beta}}$  that Make SSE as Small as Possible
  - $\widehat{\beta}_0$  and  $\widehat{\boldsymbol{\beta}}$  are Easily Found Using Matrix Representation
- Regularized Estimation
  - Produces Biased Estimates
  - Shrinks Coefficients Toward 0
  - Favors Smaller Models
  - May Lead to a Better Model for Out-of-Sample Prediction

## Part 2: Shrinkage Estimation and More Meditation



- Three Popular Methods
  - Download R Package

```
> library(glmnet)
```

- Penalized SSE

$$PSSE = SSE + \lambda[(1 - \alpha) \sum_{i=1}^p \beta_i^2 + \alpha \sum_{i=1}^p |\beta_i|]$$

- Variations

- Ridge (1970):  $\lambda = 1$  &  $\alpha = 0$
- Lasso (1996):  $\lambda = 1$  &  $\alpha = 1$
- Elastic Net (2005)  
 $\lambda = 1$  &  $0 < \alpha < 1$


- Notice When

- $\lambda = 0 \Rightarrow PSSE = SSE$
- As  $\lambda$  Gets Bigger, the Coefficients Approach 0

## Part 2: Shrinkage Estimation and More Meditation

- Run Chunk 1
  - Ridge Penalty


```
> ridge.mod=glmnet(x=as.matrix(SIM.DATA[,-1]),  
+                  y=as.vector(SIM.DATA[,1]),  
+                  alpha=0)  
> plot(ridge.mod,xvar="lambda")
```



## Part 2: Shrinkage Estimation and More Meditation

- Chunk 1 (Continued)
  - Lasso Penalty


```
> lasso.mod=glmnet(x=as.matrix(SIM.DATA[,-1]),  
+                  y=as.vector(SIM.DATA[,1]),  
+                  alpha=1)  
> plot(lasso.mod,xvar="lambda")
```



## Part 2: Shrinkage Estimation and More Meditation

- Chunk 1 (Continued)
  - Elastic Net Penalty

```
> enet.mod=glmnet(x=as.matrix(SIM.DATA[,-1]),  
+                  y=as.vector(SIM.DATA[,1]),  
+                  alpha=1/2)  
> plot(enet.mod,xvar="lambda")
```





Irrelevant  
Nonsense



Watch Me  
Whip  
Watch Me  
Lasso

## Part 2: Shrinkage Estimation and More Meditation



- Tuning Parameters
  - Use Cross-Validation to Choose Tuning Parameters  $\lambda$  &  $\alpha$
  - Constraints
    - $\lambda > 0$
    - $0 \leq \alpha \leq 1$
  - Best Approach:
    - Divide Data Into Train & Test
    - Loop Over a Vector of Alpha
    - Find Best Lambda for Each Alpha Considered Using CV in Train
    - For Each Alpha and Best Lambda, Predict on Test and Select Alpha and Lambda that Minimize MSE

## Part 2: Shrinkage Estimation and More Meditation



- Chunk 2
  - Illustration of 10 Fold CV
  - Finding Best Combination of Alpha and Lambda

	alpha	lambda	MSE
[1,]	0.0	25.073407	186.1875
[2,]	0.1	8.601355	149.4262
[3,]	0.2	4.719988	137.4652
[4,]	0.3	3.453454	133.7793
[5,]	0.4	2.590091	130.7873
[6,]	0.5	2.274097	130.5494
[7,]	0.6	1.895081	129.1495
[8,]	0.7	1.782728	130.0287
[9,]	0.8	1.559887	129.2083
[10,]	0.9	1.521754	131.2262
[11,]	1.0	1.247909	128.0857



Best:  $\lambda = 1$  &  $\alpha = 1.24$

## Part 3: Less Meditation and More Application



- Built-In Data `> mpg`
  - $n=234$
  - Focus is on Modeling Hwy MPG
  - Subset Data to Include Only Wanted Covariates

year <int>	displ <dbl>	cyl <int>	drv <chr>	cty <int>	hwy <int>	fl <chr>	class <chr>
1999	1.8	4	f	18	29	p	compact
1999	1.8	4	f	21	29	p	compact
2008	2.0	4	f	20	31	p	compact
2008	2.0	4	f	21	30	p	compact
1999	2.8	6	f	16	26	p	compact
1999	2.8	6	f	18	26	p	compact

- There are  $p=7$  Covariates
- Difficulty
  - Fitting all Combinations
  - Considering All 2-Way Interaction Terms

## Part 3: Less Meditation and More Application



- Run Chunk 1
  - Creating Model Matrix
    - Up to 2-Way Interactions
    - Now,  $p=115$
  - Model Selection is Difficult
  - Dividing Data into Train & Test is Not Advised ( $n=234$ )
- Run Chunk 2
  - Only a Few Options

<b>alpha</b> <dbl>	<b>lambda</b> <dbl>	<b>CV.Error</b> <dbl>
0.00	1.44063441	1.722966
0.25	0.55006214	1.620769
0.50	0.18956825	1.488094
0.75	0.10492193	1.456773
1.00	0.04942052	1.411025

Lowest Estimation of Prediction Error

## Part 3: Less Meditation and More Application



- Chunk 2 (Continued)
  - Understanding cv.glmnet Object
    - `$lambda` = Contains Vector of Lambda Auto-Generated
    - `$cvm` = Cross Validated Estimate of Error for Each Lambda in `$lambda`
    - `$lambda.min` = The Lambda that Leads to Smallest CV Measure of Error
    - `$lambda.1se` = The Largest Value of Lambda Such That Error is Within 1 SD of the Error Using `$lambda.min`

## Part 3: Less Meditation and More Application



- Run Chunk 3
  - Next
    - Use Best Alpha and Lambda
    - Observe the Non-Zero Coefficients
    - Plot Predictions and Errors


- Table of Non-Zero Coefficients
  - Before  $p=115$
  - Now  $p=28$

```
## # A tibble: 29 x 2
##   Parameter      Estimate
##   <chr>          <dbl>
## 1 Int            -123.
## 2 year           0.0660
## 3 cty            0.799
## 4 flr           -1.37
## 5 flr           -0.0629
## 6 classpickup   -0.104
## 7 classsuv      -1.37
## 8 year:cyl      -0.0000392
## 9 year:drv      0.0000955
## 10 year:cty     0.0000565
## 11 year:classmidsize 0.0000259
## 12 year:classpickup -0.000659
## 13 displ:drv     0.127
## 14 displ:classmidsize 0.0317
## 15 displ:classsuv -0.178
## 16 cyl:flr       -0.143
## 17 cyl:flr       -0.0973
## 18 cyl:classcompact 0.0462
## 19 cyl:classsuv  -0.0262
## 20 drv:cty       0.0466
## 21 drv:cty       0.0282
## 22 drv:fld       2.54
## 23 drv:classsubcompact -0.0754
## 24 cty:classminivan -0.0574
## 25 cty:classpickup -0.106
## 26 flr:classmidsize 0.488
## 27 flr:classsubcompact -1.42
## 28 fld:classsuv  -0.552
## 29 flr:classsuv  -0.431
```

## Part 3: Less Meditation and More Application



- Chunk 3 (Continued)
  - Comparing Predict and Actual






## Part 3: Less Meditation and More Application



- Chunk 3 (Continued)
  - Distribution of Residuals



Closing



Disperse  
and Make  
Reasonable  
Decisions